

# Idle Scheduling in Linux

Nikhil Rao  
Google, Inc.

# Outline

- Workloads and CPU QoS goals
- Lessons learnt
- Limitations in balancing low weight task groups
- Solutions
- Future work

# Types of Workloads

## **Latency sensitive applications**

- High priority workload
- Strict latency, throughput guarantees
- Typical request-response tasks
- User-facing services like search, maps, etc.

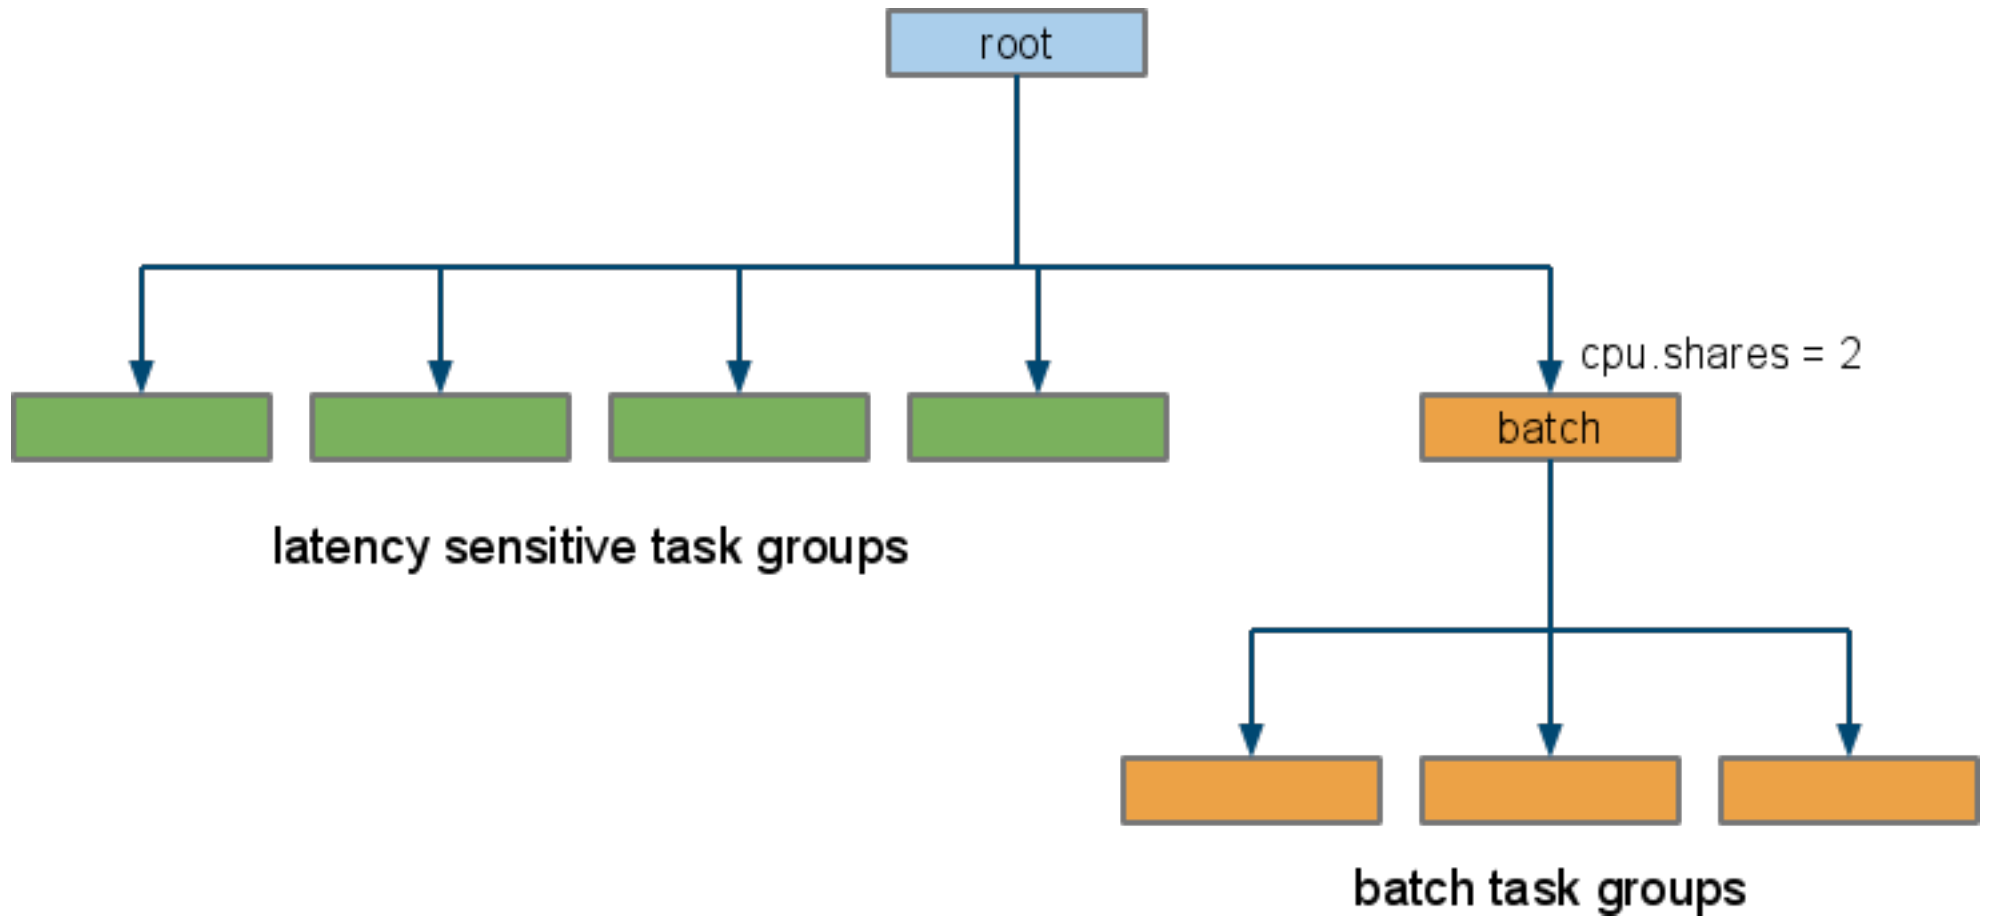
## **Batch applications**

- Low priority workload
- Soak up idle resources on the machine
- Usually cpu soakers with little or no I/O
- Long running batch jobs like video transcoding, etc.

# QoS Requirements

- Isolation between latency sensitive applications
- Guaranteed latency response, fairness for latency sensitive applications
- Strict priority for latency sensitive over batch
- Maximize utilization if there is demand for cpu

# Task group structure



# Lessons Learnt

## Wins!

- Group scheduling works well!
  - Good isolation, fairness between latency sensitive applications
- Shares is simple abstraction for application developers
- Good isolation between latency sensitive and batch tasks
  - Could be improved, future work

## Pain Point:

- Starvation and degraded end-to-end latency reported for batch tasks!
  - Large weight differentials lead to sub-optimal utilization
  - Insufficient resolution for balancing low weight groups

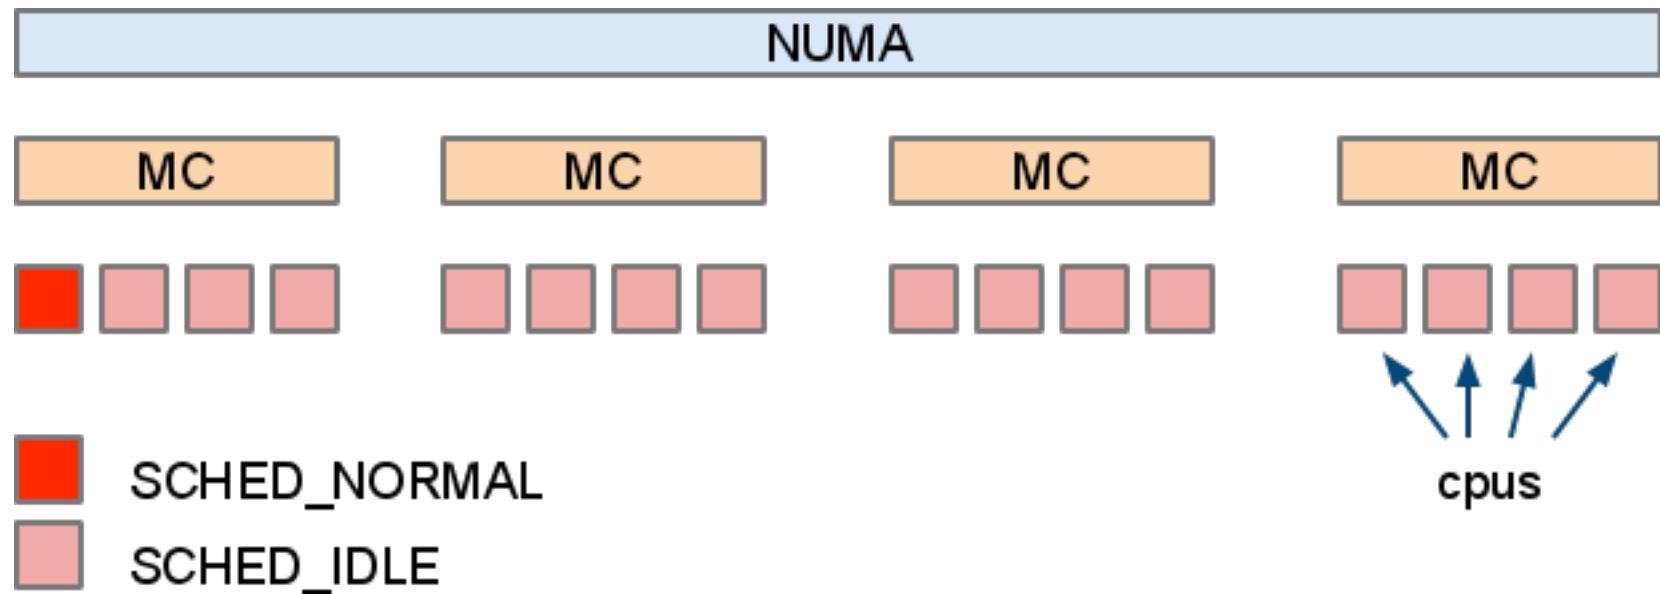
# #1: Sub-optimal Utilization Example

## Test Setup

- 16 cpu test machine (quad-socket, quad-core)
- 15 SCHED\_IDLE soaker tasks (load wt = 3)
- 1 SCHED\_NORMAL soaker task (nice 0, load wt = 1024)
- 2.6.36 kernel

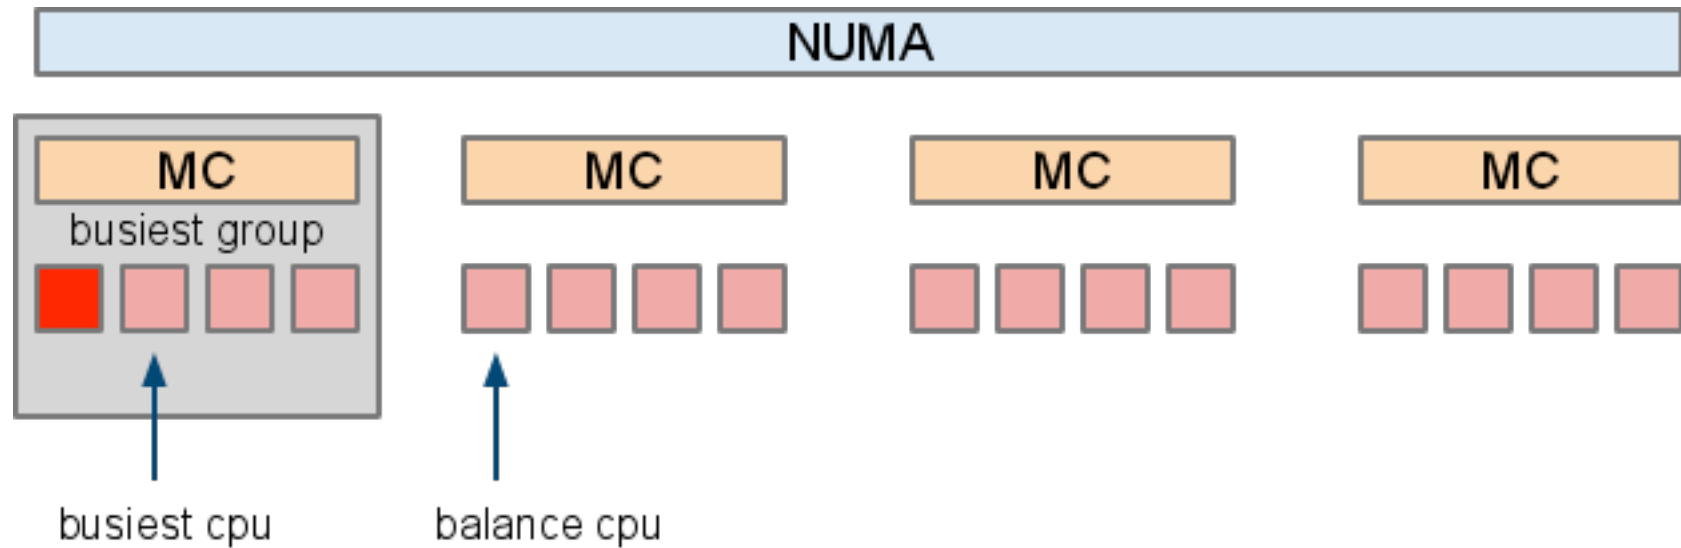
## Result

04:58:46 PM	CPU	%user	%nice	%sys	%iowait	%irq	%soft	%steal	%idle	intr/s
04:58:47 PM	all	81.47	0.00	0.25	0.00	0.00	0.00	18.28	13796.00	
04:58:48 PM	all	81.20	0.00	0.25	0.00	0.00	0.00	18.55	13816.00	
04:58:49 PM	all	80.93	0.19	0.25	0.00	0.00	0.06	18.57	13965.00	
04:58:50 PM	all	81.40	0.00	0.25	0.00	0.00	0.00	18.35	13837.37	
04:58:51 PM	all	81.19	0.00	0.31	0.00	0.00	0.00	18.50	13592.08	
04:58:52 PM	all	81.25	0.00	0.25	0.00	0.00	0.00	18.50	13721.00	
04:58:53 PM	all	81.19	0.00	0.25	0.00	0.00	0.00	18.56	13764.00	
04:58:54 PM	all	81.25	0.00	0.25	0.00	0.00	0.00	18.50	13841.41	
04:58:55 PM	all	80.30	0.00	1.19	0.00	0.00	0.00	18.51	14989.11	
04:58:56 PM	all	80.77	0.00	0.50	0.00	0.00	0.00	18.73	13964.65	
<b>Average:</b>	<b>all</b>	<b>81.09</b>	<b>0.02</b>	<b>0.37</b>	<b>0.00</b>	<b>0.00</b>	<b>0.01</b>	<b>0.00</b>	<b>18.51</b>	<b>13929.53</b>

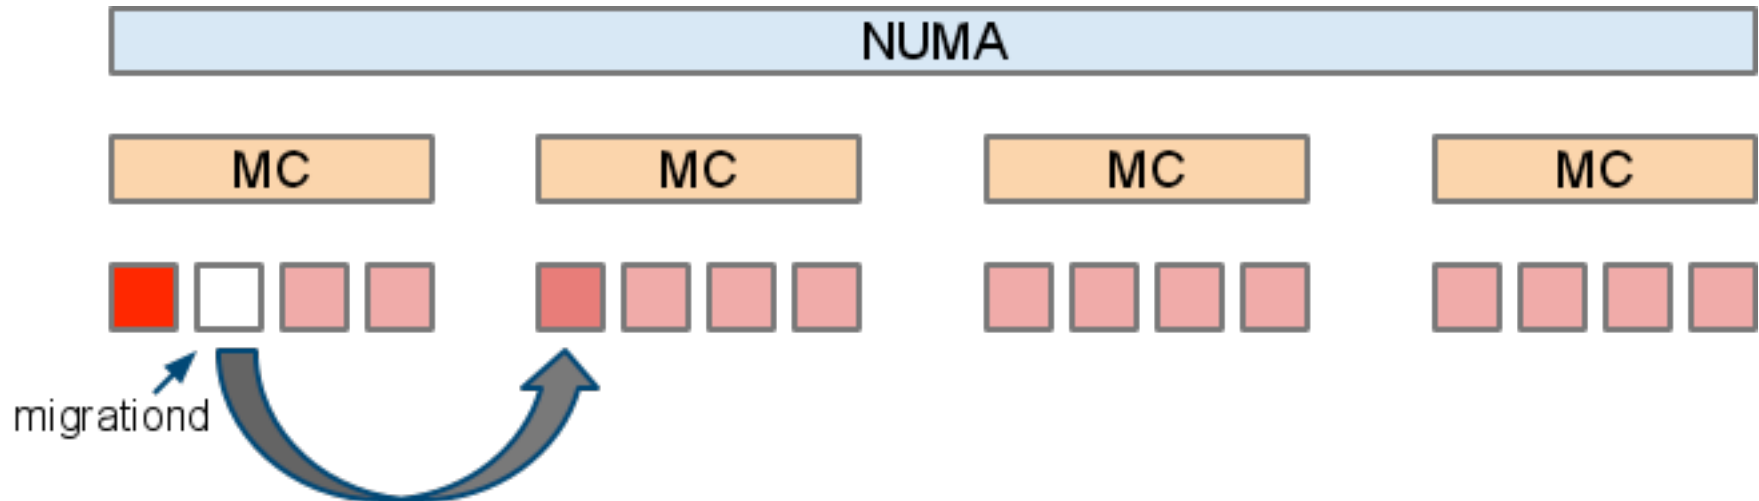


- Two scheduling domain levels, MC and NUMA
- cpu0 has the SCHED\_NORMAL task
- SCHED\_IDLE tasks distributed equally on cpus 1-15






- Consider balancing decisions at the NUMA domain
- Sched group weights = { 1033, 12, 12, 12 }
- `f_b_g()` picks group0 as busiest group
- `f_b_q()` picks cpu1, 2 or 3 as busiest queue
  - Never picks cpu0 because weight > imbalance

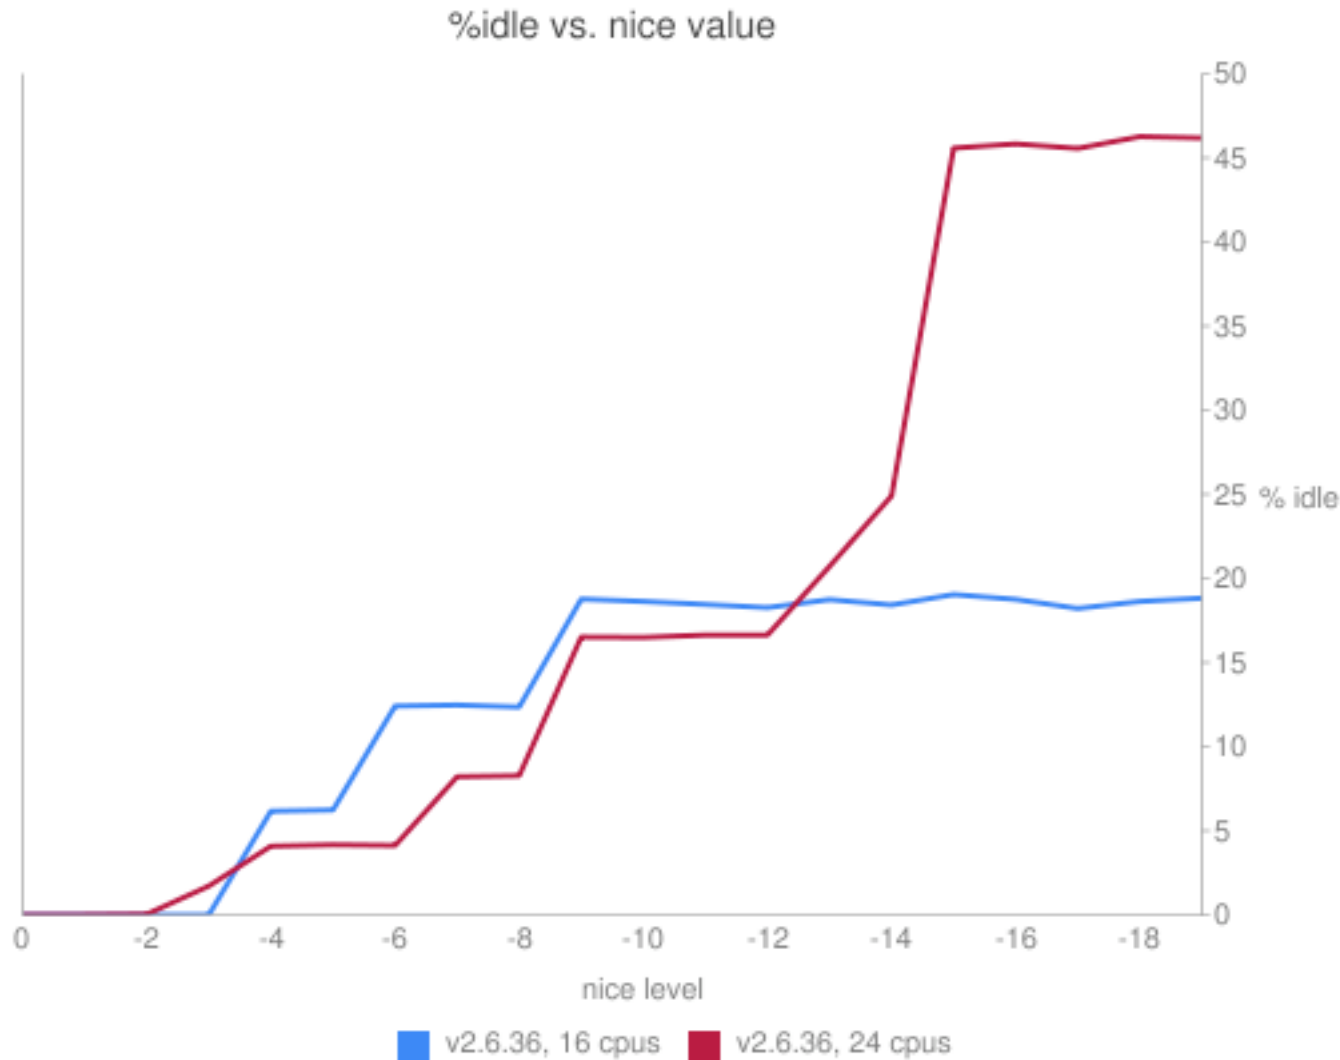


- Load balancer pulls tasks until one task per runqueue
- Balancing operations fail (unable to pull running task)
- Active migration kicks in after 5 failures
- Pushes SCHED\_IDLE task from the sched group



- Active migrations kick off all SCHED\_IDLE tasks
- Idle cpus unable to pull load back
  - load balancer does not find any busy group

# #1: Experiment with niced task



Increasing task priority (via nice) leads to sub-optimal utilization!

# #1: Load Balancer Fixes

- Introduced notion of extra group capacity
  - extra capacity  $\Rightarrow$  `nr_running < group_capacity`
- Set `group_imb` only if `max_nr_running > 1`
- Group capacity fixes when `SD_PREFER_SIBLING` enabled on child domain
- Force balancing if local group has extra capacity

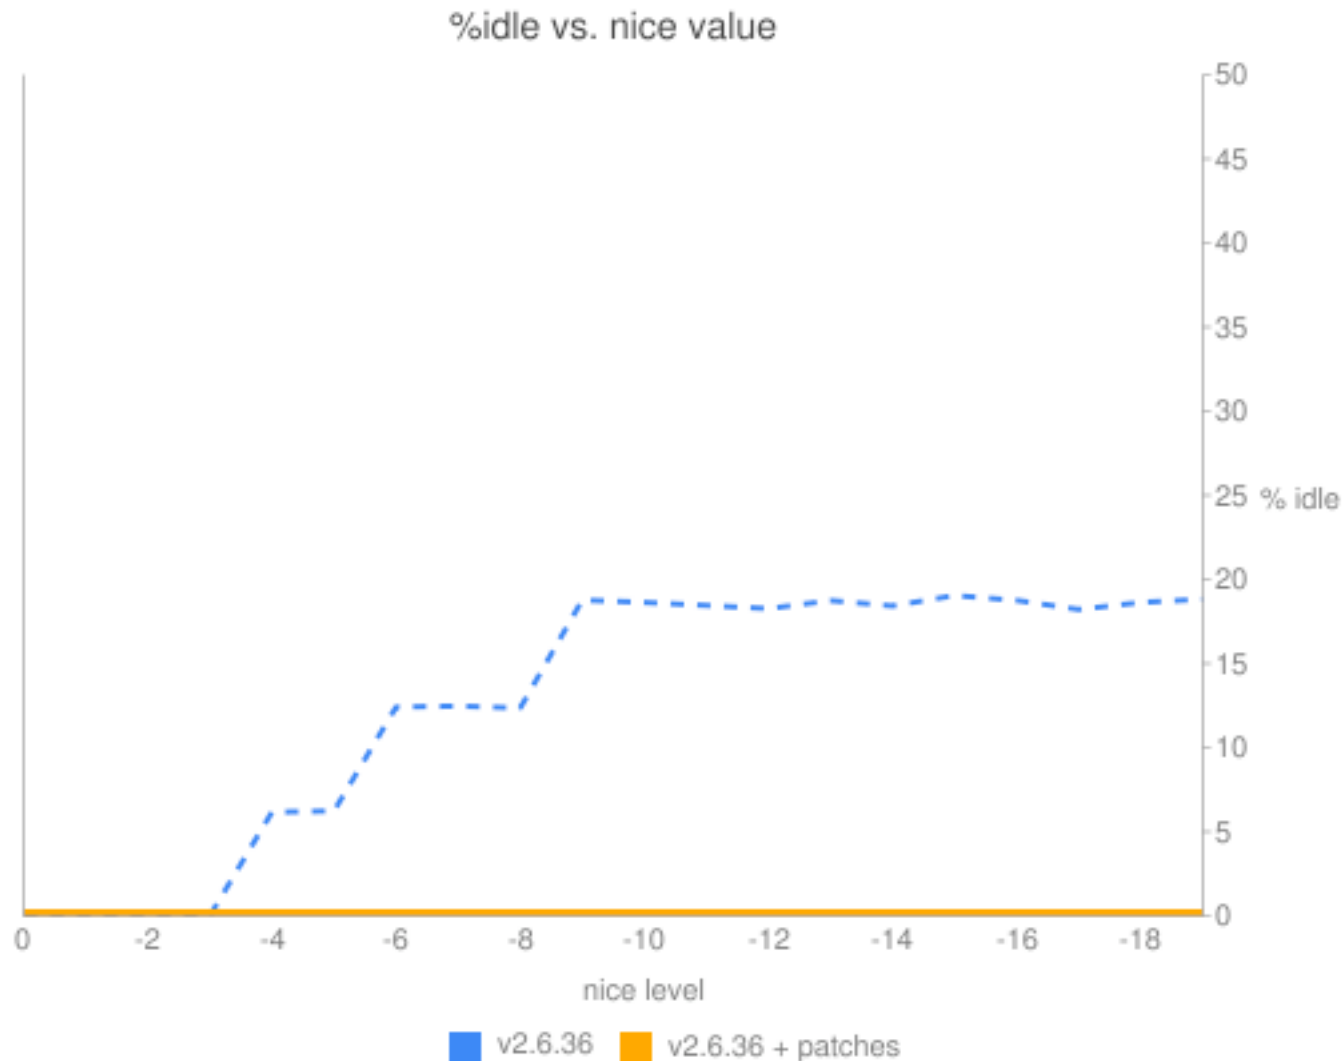
# #1: Results after fixes

- v2.6.36 + patches
- 15 SCHED\_IDLE tasks, 1 SCHED\_NORMAL task
- Improved utilization!

12:58:29 PM	CPU	%user	%nice	%sys	%iowait	%irq	%soft	%steal	%idle	intr/s
12:58:30 PM	all	99.81	0.00	0.19	0.00	0.00	0.00	0.00	16384.00	
12:58:31 PM	all	99.75	0.00	0.25	0.00	0.00	0.00	0.00	16428.00	
12:58:32 PM	all	99.81	0.00	0.19	0.00	0.00	0.00	0.00	16345.00	
12:58:33 PM	all	99.75	0.00	0.25	0.00	0.00	0.00	0.00	16383.00	
12:58:34 PM	all	99.75	0.00	0.19	0.00	0.00	0.06	0.00	16333.00	
12:58:35 PM	all	99.81	0.00	0.19	0.00	0.00	0.00	0.00	16359.00	
12:58:36 PM	all	99.75	0.00	0.25	0.00	0.00	0.00	0.00	16523.23	
12:58:37 PM	all	99.75	0.00	0.25	0.00	0.00	0.00	0.00	16352.00	
12:58:38 PM	all	98.75	0.00	1.25	0.00	0.00	0.00	0.00	17128.00	
12:58:39 PM	all	99.31	0.06	0.62	0.00	0.00	0.00	0.00	16757.00	
<b>Average:</b>	<b>all</b>	<b>99.63</b>	<b>0.01</b>	<b>0.36</b>	<b>0.00</b>	<b>0.00</b>	<b>0.01</b>	<b>0.00</b>	<b>16499.20</b>	

# #1: Results after fixes

- Experiment with niced task, 16 cpu machine



## #2: Insufficient granularity

- Load balancing math breaks down with low weight groups
- Consider hierarchal load factor calculations in `tg_load_down()` for static batch task group (`shares = 2`)
  - $h\_load = parent \rightarrow h\_load * shares / parent \rightarrow cfs\_rq \rightarrow wt + 1$
  - `h_load = 1` for `/batch` task group
  - `h_load = 0` for any task group under `/batch`
- Setting `h_load = 0` leads to a couple of issues:
  - Herd migrations
  - Loss of fairness between batch tasks



# #2: Herd Migrations

- Herd migrations: mass migration of batch tasks from the busiest cpu to the balancing cpu to satisfy large imbalance
- Insufficient granularity of low weight task groups result in h\_load being estimated as 0
- Small imbalance is greatly exaggerated
  - For example, imbalance of 10 with 5 tasks on busiest cpu results translates to rem\_move of ~51K
  - Enough to migrate all tasks except running task
- Incorrect accounting after migration
  - $\text{moved\_load} = \text{moved\_load} * \text{h\_load} / (\text{weight} + 1) = 0!$
  - Failed migration!

# #2: Example of Herd Migration

## Test Setup

- 16 cpu test machine (quad-socket, quad-core)
- Create a batch task group under /batch with shares = 1024
- Spawn 48 tasks with random sleep/busy pattern (100ms)

## Result

03:04:24 PM	CPU	%user	%nice	%sys	%iowait	%irq	%soft	%steal	%idle	intr/s
03:04:25 PM	all	91.72	0.00	0.31	0.00	0.00	0.00	7.97	15607.00	
03:04:26 PM	all	93.50	0.00	0.37	0.00	0.00	0.00	6.12	15749.50	
03:04:27 PM	all	94.62	0.00	0.31	0.00	0.00	0.00	5.06	16045.00	
03:04:28 PM	all	94.69	0.00	0.31	0.00	0.00	0.00	4.99	16311.11	
03:04:29 PM	all	93.95	0.00	0.37	0.00	0.00	0.00	5.68	16037.00	
03:04:30 PM	all	94.07	0.00	0.37	0.00	0.00	0.00	5.56	15843.56	
03:04:31 PM	all	94.93	0.00	0.31	0.00	0.00	0.00	4.75	16081.00	
03:04:32 PM	all	95.19	0.00	0.38	0.00	0.00	0.00	4.44	16157.00	
03:04:33 PM	all	95.75	0.00	0.37	0.00	0.00	0.00	3.87	16030.69	
03:04:34 PM	all	95.69	0.00	0.31	0.00	0.00	0.00	4.00	16184.00	
<b>Average:</b>	<b>all</b>	<b>94.41</b>	<b>0.00</b>	<b>0.34</b>	<b>0.00</b>	<b>0.00</b>	<b>0.00</b>	<b>5.25</b>	<b>16003.89</b>	

# #2: Example of Herd Migration



# #2: Lack of Fairness (batch tasks)

## Test Setup:


- 16 cpu machine (quad-core, quad-socket)
- 48 task groups with one soaker each
- Compare `cpuacct.usage` for each task group

## Result:

Distribution of runtime between batch tasks



Distribution of runtime between latency sensitive tasks



## #2: Ideas to improve granularity

- Scale up shares by a constant
  - Update MIN\_SHARES, MAX\_SHARES
  - Update nice to weight ratios
  - Loss of accuracy in update\_curr()
    - Scale down weights before calling update\_curr()
    - Can we do 128-bit math?
- Scale load weights before balancing operations
  - Scale h\_load by a factor of 1024 in load\_balance\_fair()

# #2: Results of scaling up shares

- Reduces herd migrations, improves utilization!

02:43:04 PM	CPU	%user	%nice	%sys	%iowait	%irq	%soft	%steal	%idle	intr/s
02:43:05 PM	all	99.56	0.00	0.31	0.00	0.00	0.00	0.12	16772.00	
02:43:06 PM	all	98.94	0.00	1.00	0.00	0.00	0.00	0.06	17031.00	
02:43:07 PM	all	98.94	0.00	0.75	0.00	0.00	0.06	0.25	17002.97	
02:43:08 PM	all	98.81	0.00	1.00	0.00	0.00	0.00	0.19	16930.30	
02:43:09 PM	all	98.75	0.00	1.00	0.00	0.00	0.00	0.25	16792.00	
02:43:10 PM	all	99.56	0.00	0.31	0.00	0.00	0.00	0.12	16785.00	
02:43:11 PM	all	99.44	0.00	0.37	0.00	0.00	0.00	0.19	16923.00	
02:43:12 PM	all	98.69	0.00	0.37	0.00	0.00	0.06	0.87	16806.00	
02:43:13 PM	all	99.25	0.00	0.37	0.00	0.00	0.00	0.37	16760.00	
02:43:14 PM	all	99.50	0.00	0.37	0.00	0.00	0.00	0.12	16515.84	
<b>Average:</b>	<b>all</b>	<b>99.14</b>	<b>0.00</b>	<b>0.59</b>	<b>0.00</b>	<b>0.00</b>	<b>0.01</b>	<b>0.00</b>	<b>0.26</b>	<b>16831.57</b>

- Improved fairness

fairness distribution of  
batch tasks with scaled  
shares



# Future Work

- Avoid batch task preempting latency sensitive task
  - Extend "SCHED\_IDLE" concept to group entities
  - Change preemption model for group entities
  - Tasks in to idle groups do not preempt non-idle tasks
- Support more complex task group hierarchies
  - Arbitrary nesting of task groups
  - Guarantee fairness to tasks in this hierarchy
  - Maximize overall system utilization
- Further improve fairness between batch tasks

# Thank you!

Nikhil Rao  
ncrao@google.com