

# Linux Plumbers Conference 2010

---

## **TCP-NV** **Congestion Avoidance for** **Data Centers**

**Lawrence Brakmo**

**Google**




# TCP Congestion Control

---

- **Algorithm for utilizing available bandwidth without too many losses**
  - No attempt at eliminating losses
  - Works well at eliminating congestion collapse
- **TCP's congestion control will continually increase its congestion window (# packets in transit) until losses occur**
  - $\text{rate} = \text{cwnd} / \text{RTT}$
- **I.e. congestion control repeatedly creates congestion and packet losses**
- **Unfairness**
  - Between flows with different RTTs (up to 20-40x)
- **Works well for large transfers and it has proven to work well in many different environments**



# Congestion Window Graph



# Data Center Perspective

---

- **RPCs make large percent of traffic**
  - Important to test TCP behavior with RPCs
- **Latency sensitive**
- **Congestion control results in queue build-up at host, switches and routers**
  - Increasing average latency of smaller size RPCs
- **Congestion control results in periodic packet losses**
  - Increasing high percentile latency of smaller size RPCs
- **Need Congestion avoidance**
  - Reduce queue buildup and losses



# Congestion Avoidance

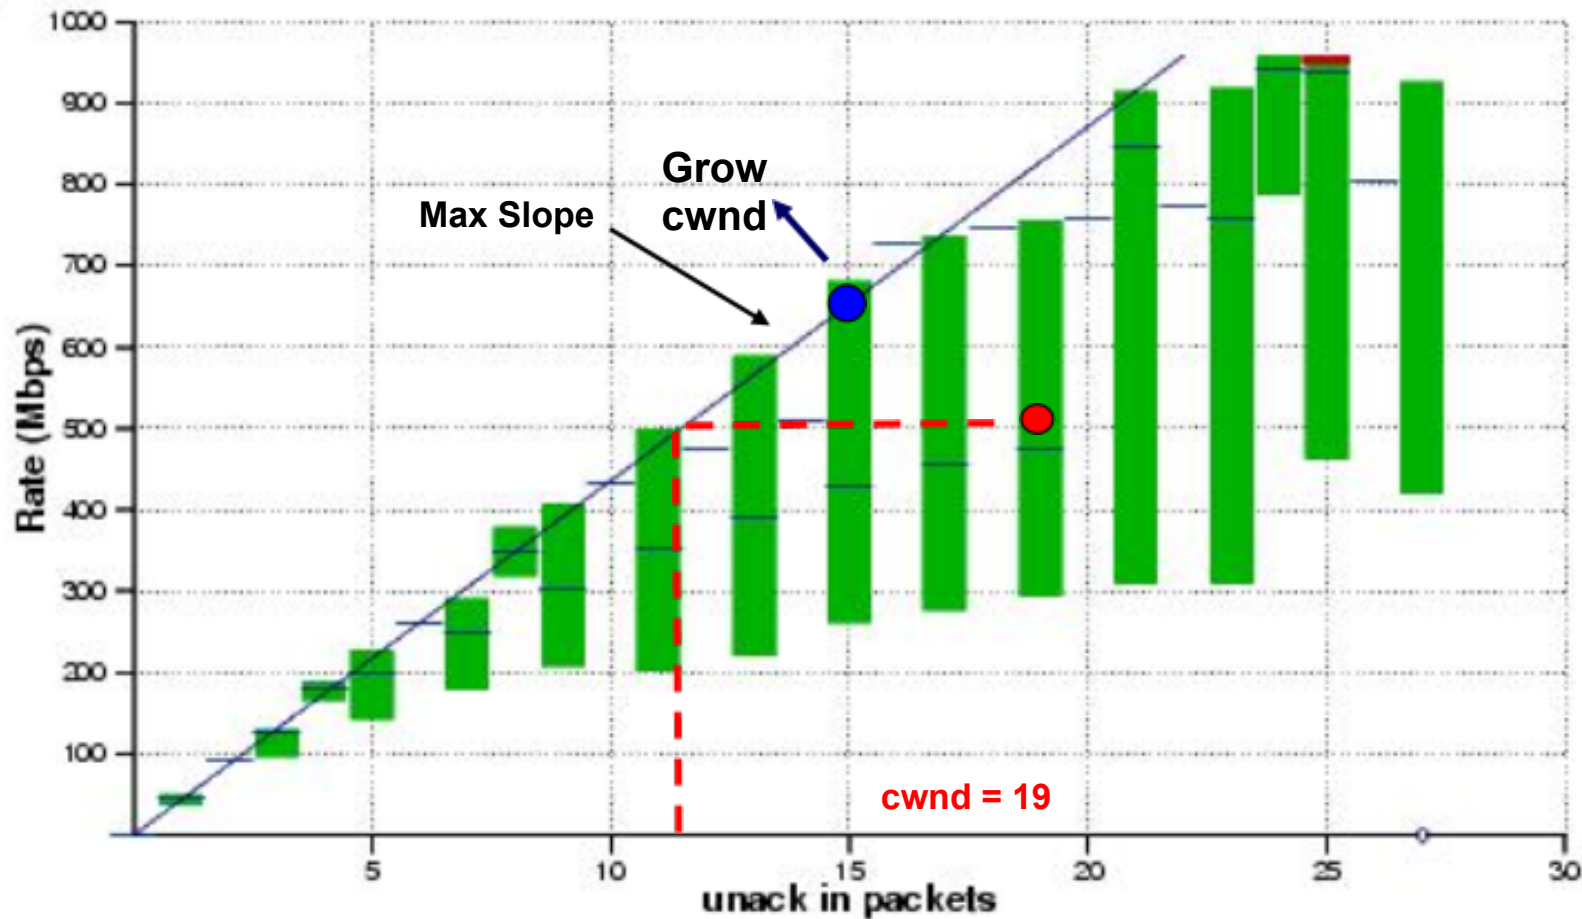
---

- **Goal: Prevent queue build-up and packet losses while making full utilization of available bandwidth**
  - Only grow congestion window when there is bandwidth available
  - Decrease congestion window when congestion (queue build-up) is detected
    - Unlike congestion control which only decreases with losses
- **Typically delay (RTT) based**
  - Increase in RTT implies queue buildup
  - Problem: RTT is VERY noisy



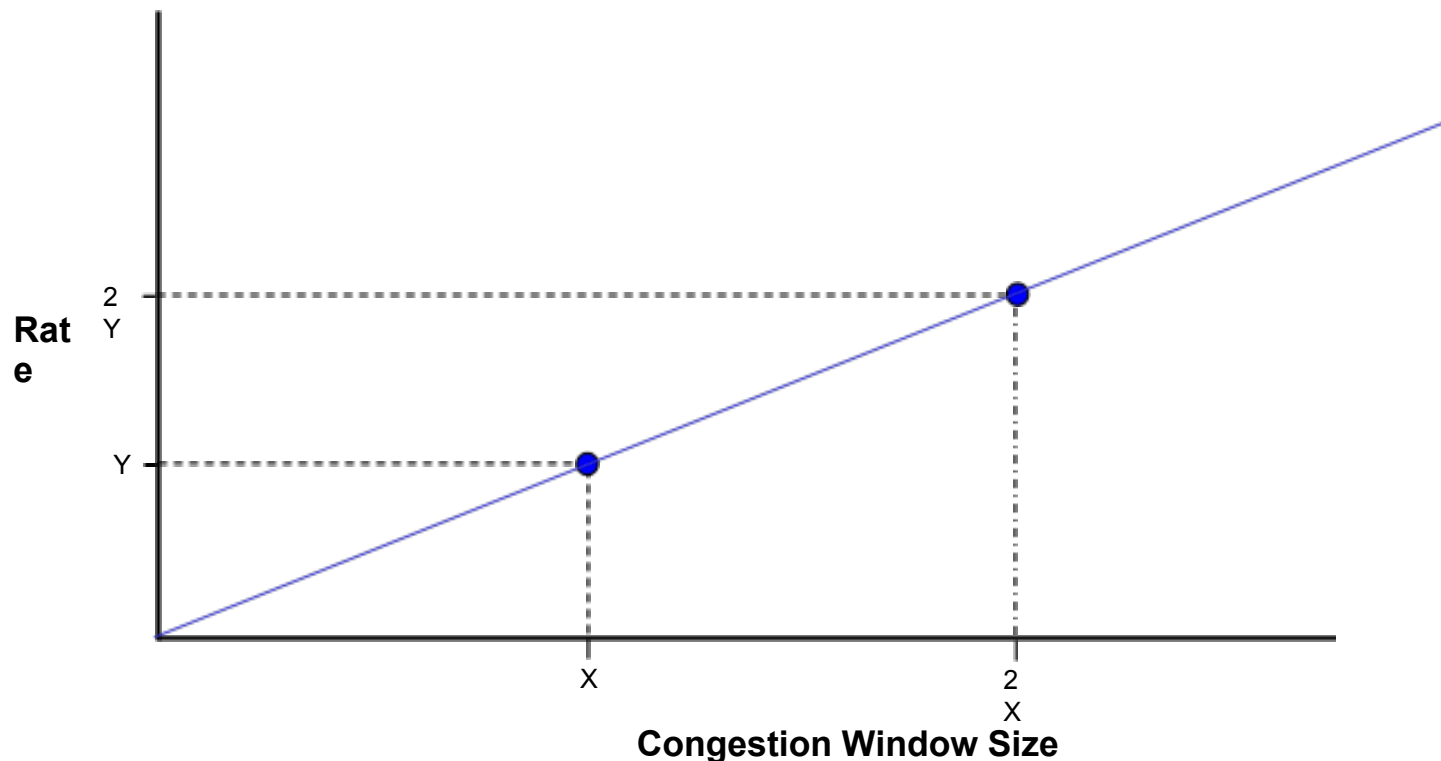
# TCP-NV Congestion Avoidance

- We can only detect the level of congestion, not the available bandwidth



# Basic Concept of TCP-NV Congestion Avoidance

- In a perfect world, and while there is no congestion, if a cwnd of  $X$  achieves rate  $Y$ , then a cwnd of  $2X$  should achieve rate  $2Y$



# TCP-NV Implementation

---

- **In Linux 2.6.34, soon in 2.6.36**
- **On top of BIC**
  - BIC behavior when there is no congestion
- **Most of the changes are self contained in one CA function: `nvtcp_acked()`**
  - Collects appropriate info
  - Every so often makes a CA decision
    - Allow BIC to grow cwnd
    - Don't allow BIC to grow cwnd
    - Decrease cwnd and don't allow BIC to grow cwnd
- **Need to store in-flight every time a packet is sent** ☐
  - Needed by `ntcp_acked()`, stored in `skb`
- **Can be disabled through `sysctl` (in which case it behaves like BIC)**






## More Implementation Details

---

- **Every time a packet is sent**
  - Store time in us, bytes in-flight
- **Every time an ACK is received**
  - Use sent time and in-flight for newest packet ACKed
  - Calculate slope = rate/in-flight
  - Keep largest slope seen so far (recalc every so often)
  - Keep largest rate seen since last CA decision
- **Every so often**
  - $\text{pred\_cwnd} = \text{max\_rate} / \text{max\_slope}$
  - if  $\text{pred\_cwnd} + \text{pad} > \text{cwnd}$ 
    - Reduce cwnd by percent of congestion
    - Stop BIC from growing cwnd
  - else
    - Allow BIC to grow cwnd



# Congestion Window Graph



# Issues

---

- **TSO, LRO and interrupt coalescence**
  - Reduces number of data points
    - May only get one ACK per RTT
  - Need to inflate window for one flow to use full bandwidth
- **Coexistence with congestion control**
  - As CC increases cwnd, CA reduces its own
  - Usually cannot do both
  - Unless RTT of CA is  $\ll$  RTT of CC
    - Traffic within DC CA, traffic outside DC CC
- **TCP-NV needs pacing when RTT > 1 to 5ms**
  - due to big bursts when starting a new RPC
  - not an issue for traffic within a data center or cluster
- **Cannot do CA with small RPCs**
  - Needs a couple of RTTs of data



# CC vs CA Fairness

		NV	
		small RTT	large RTT
BIC	small RTT	BIC 5X > NV	BIC 10-100X > NV (avoid)
	large RTT	NV 2-5X > BIC (= BIC vs. BIC) NV losses ~0	Fair

## Issues ...

---

- **Reverse congestion**

- Congestion in the ACK direction will increase RTT
- CA mechanism will react to this

- **Solutions** □

- Prioritize pure ACKs in host and switches/routers
  - so they are not affected by congestion
- Measure reverse delay and adjust appropriately
  - New TCP option
  - Use TCP us Timestamp option
- Limit how much NV can decrease the cwnd



## Issues ...

---

- **Too many flows decrease effectiveness of Congestion Avoidance**
  - # of flows w/o losses is a function of router buffer size
  - With 1.2MB buffer size, can handle up to 128 bulk flows with no losses
    - This has improved and we are working on further improvements



# Results

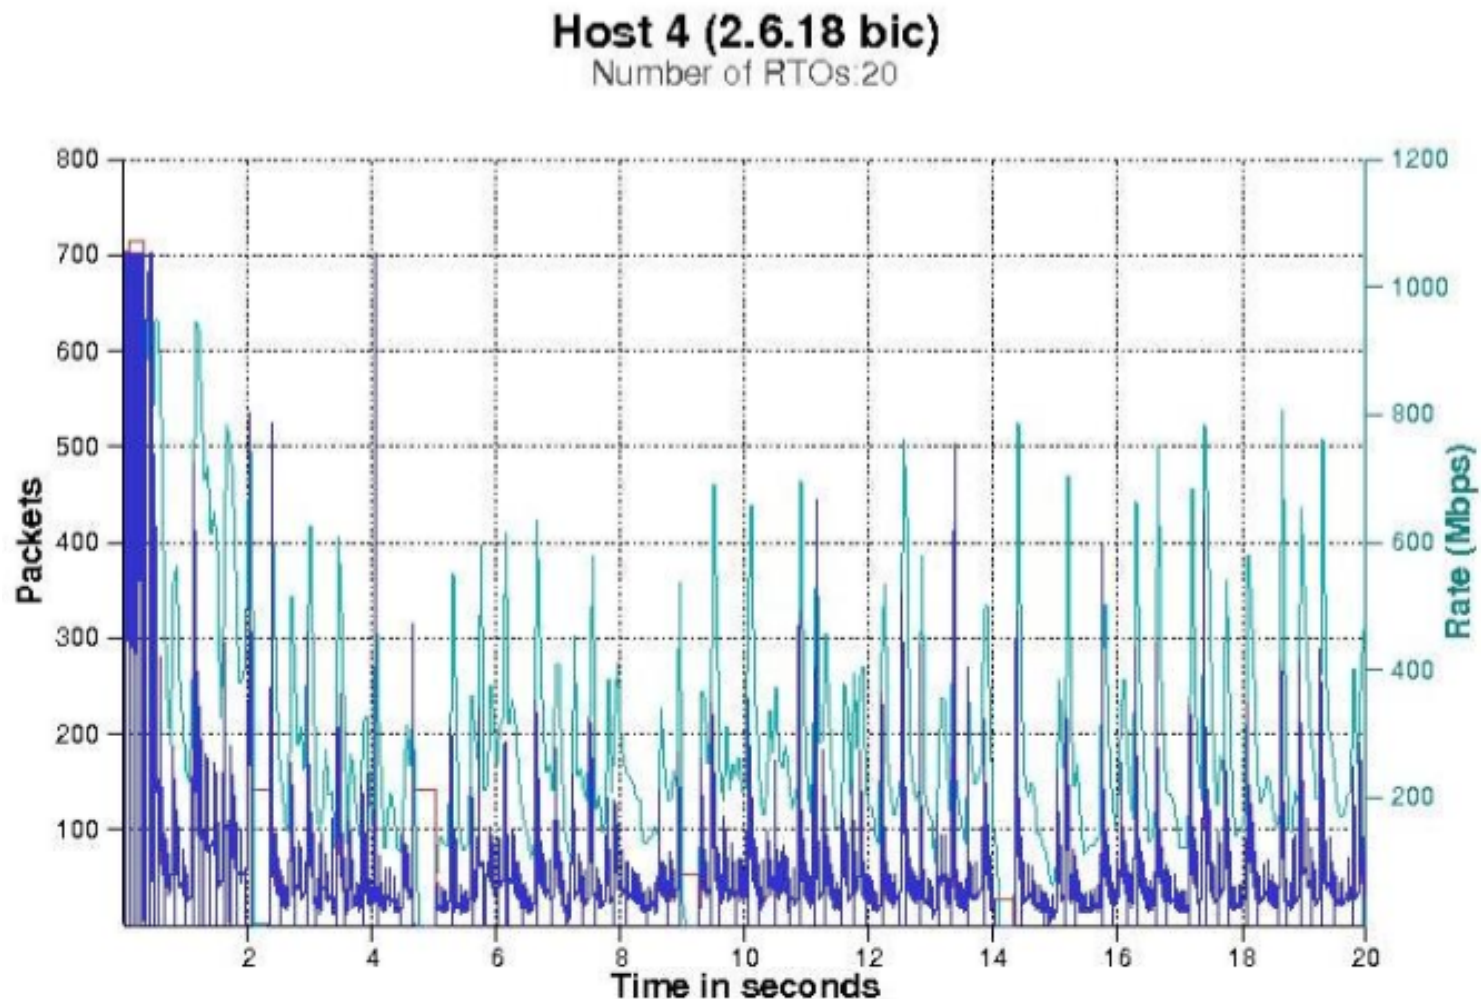
---

- **Simulations using actual Linux network stack**
  - Using Sam Jansen's Network Simulation Cradle
  - Up to 1000 hosts, 50us to 100ms RTTs
  - Can examine detail behavior of protocol, routers, etc.
  - Can easily move the code to Linux
- **Rack tests with 1G and 10G NICs**
  - Actual hardware
    - Effects of TSO, LRO, interrupt coalescence
- **Starting tests with production workloads**
  - More machines (currently 38)
  - Actual workloads




# TCP Burstiness with Small RTT Large RPCs

- Experiment: 128 flows, 100us RTT
- Graph: In-flight (dark blue) and rate (aqua) of 1 flow
- Each large peak is the result of starting a new RPC.






# TCP Burstiness with TCP-NV



# Router Queues with 128 flows

1-8 MB RPCs, 100us RTT, 1.2MB switch buffer

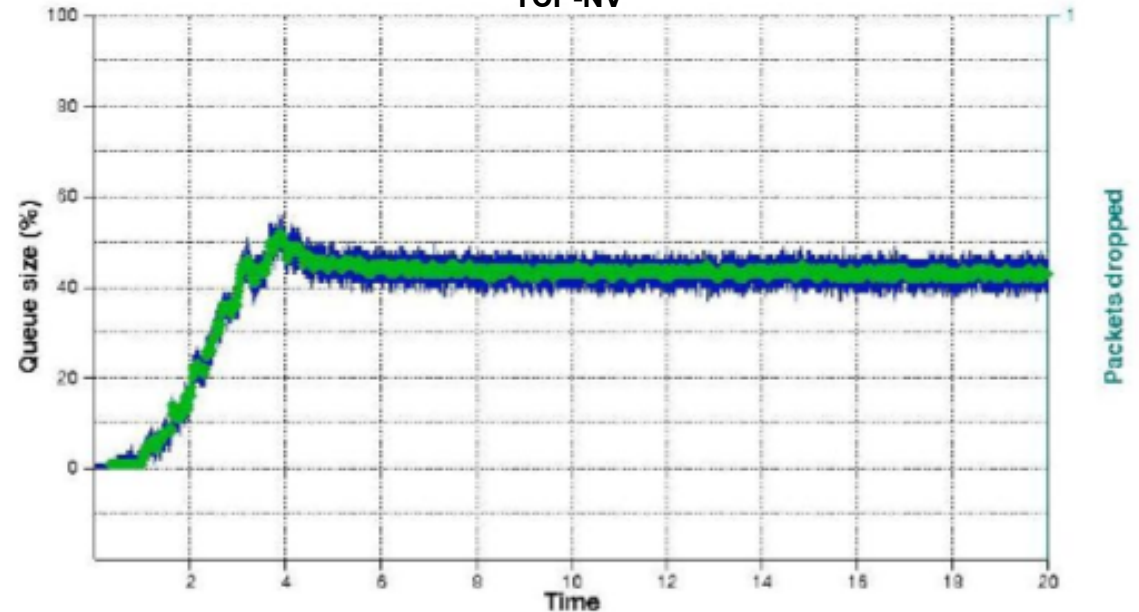
TCP-BIC



13,000 packet losses/s

0 packet drops


TCP-NV



# Router Queues with 128 flows


## 1-8 MB RPCs, 100us RTT, 175KB switch buffer

TCP-BIC



17,800 packet drops/s

TCP-NV



75 packet drops/s




# Simulator, 10Gbps, ~1MB RPCs

	<b>TCP-BIC avg of 2,8,32 flows</b>	<b>TCP-NV avg of 2,8,32 flows</b>	<b>TCP-BIC 128 flows</b>	<b>TCP-NV 128 flows</b>
<b>Goodput (Gbps)</b>	8.1	8.8	9.5	9.6
<b>Fairness</b>	0.96	0.94	0.82	0.93
<b>Average latency</b>	642 us	251 us	1.7 ms	527 us
<b>90% latency</b>	448 us	274 us	589 us	559 us
<b>95% latency</b>	494 us	285 us	14.5 ms	584 us
<b>99% latency</b>	11.9 ms	305 us	21.0 ms	616 us
<b>Average queue size</b>	94.6	39.3	152	194
<b>Packets dropped/sec</b>	19000	0	63400	5300

## 2 1G hosts sending to 1G host (Rack test)

	Flows per host	Busy %	avg cwnd	avg RTT (ms)	Aggregate Rate (Mbps)	Losses %
<b>BIC</b>	1	9.3	119	2.2	918	1.500
<b>NV</b>	1	3.3	14	1.0	925	0.000
<b>BIC</b>	4	11.0	25	1.9	952	2.800
<b>NV</b>	4	6.0	6.4	1.0	944	0.001
<b>BIC</b>	8	9.6	22	1.8	981	3.800
<b>NV</b>	8	6.5	6.5	1.0	974	0.008





## 2 1G hosts sending to 1G host, Flow Control Enabled

	Flows per host	Busy %	avg cwnd	avg RTT (ms)	Aggregate Rate (Mbps)
<b>BIC</b>	1	7.5	737	15	943
<b>NV</b>	1	2.9	14	1	938
<b>BIC</b>	4	8.7	736	66	944
<b>NV</b>	4	5.9	7	1	944
<b>BIC</b>	8	8.9	737	134	944
<b>NV</b>	8	7.2	8	1.5	944



## Distributed Sorting (38 1Gbps machines, real HW)

TCP version	Execution Time	Retransmits	RTOs
Cubic	baseline	baseline	baseline
BIC	-3%	+36%	-25
NV	-5%	-45%	-17%



# Conclusions

---

- **Congestion avoidance is possible in the data center**
  - Can insure everyone is running CA
- **Can reduce losses and queue buildup**
  - But these are workload dependent
- **Can reduce cpu utilization**
  - Due to fewer losses and smaller host queues
- **Small changes to the network stack**
  - Most of the code in a new ca module.
- **Started running larger tests**
  - multi-rack to cluster size
- **Code will be released in December**





The End